

Mini project 3 report: Memory virtualization #2

Titsanu	Wantaratorn	57070501015
Pongtorn	Siripraparwan	57070501028
Wiroat	Saeheng	57070501039
Chutichai	Tuntithawatchaikul	57070501066
Rachasak	Ragkamnerd	57070501075



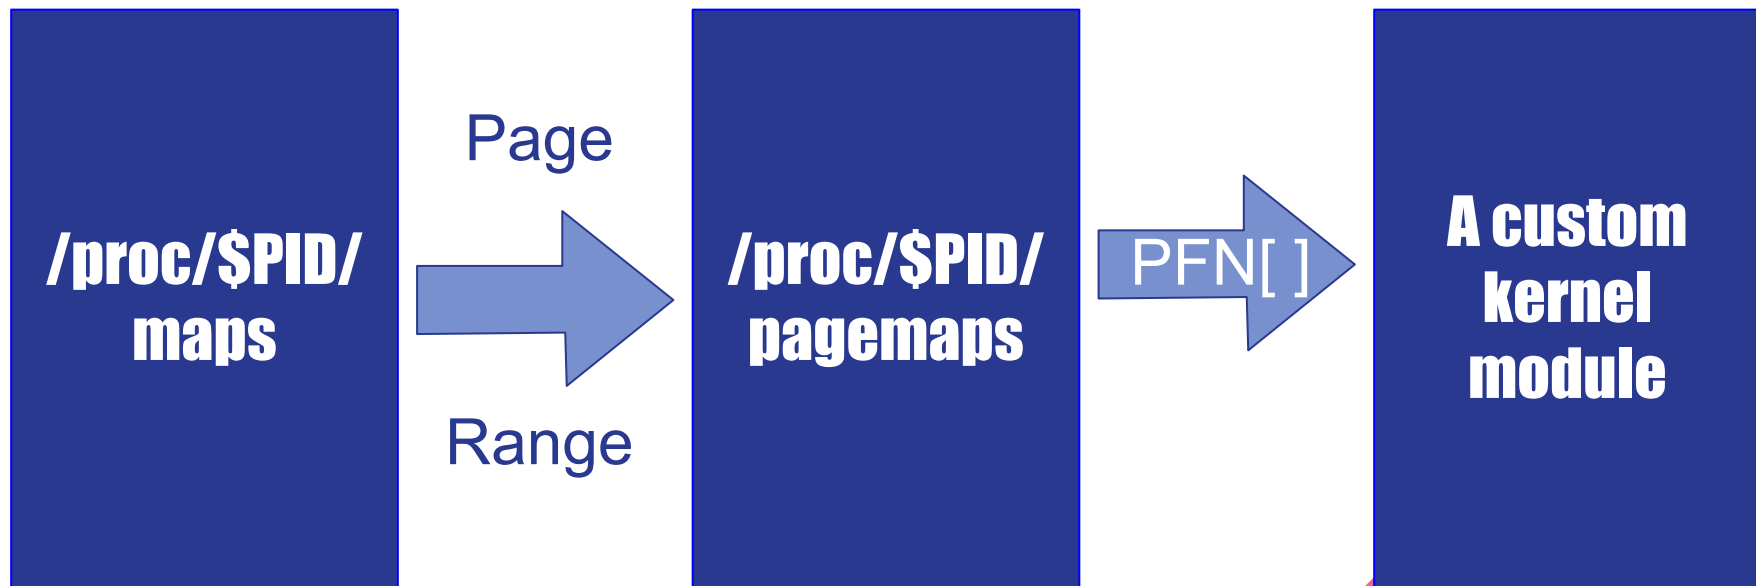
Overview

- **How it work**
 - **What we have done**
 - Reading `/proc/$PID/maps`
 - Calculating and gathering **PFN**
 - Gathering **Page flag**
 - **Problem**
 - **Demo**
 - **Q&A**
-

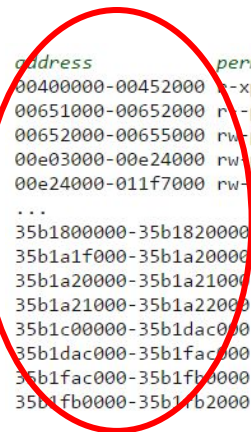


How it work

What we have done



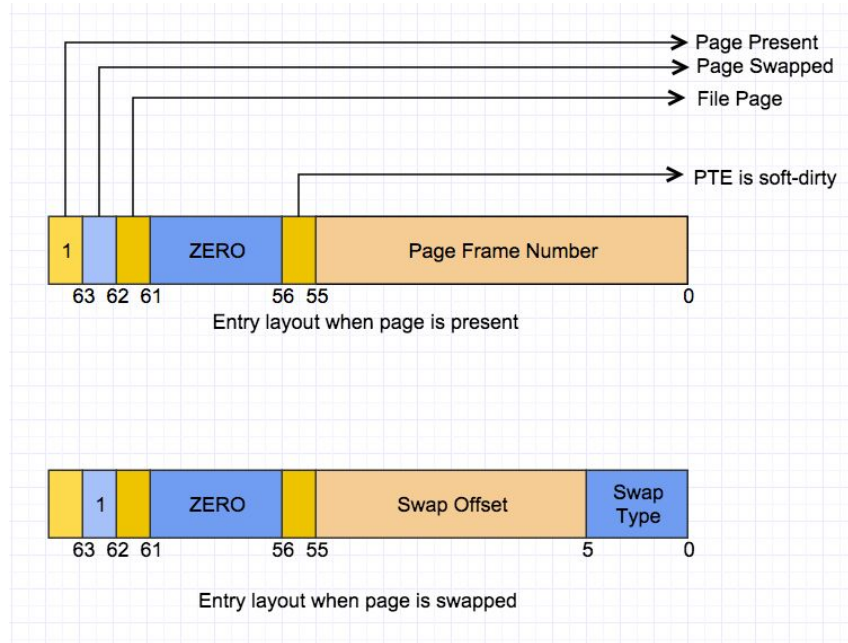
/proc/\$PID/maps



```
address      perms offset  dev   inode      pathname
00400000-00452000 r-xp 00000000 08:02 173521    /usr/bin/dbus-daemon
00651000-00652000 r--p 00051000 08:02 173521    /usr/bin/dbus-daemon
00652000-00655000 rw-p 00052000 08:02 173521    /usr/bin/dbus-daemon
00e03000-00e24000 rw-p 00000000 00:00 0        [heap]
00e24000-011f7000 rw-p 00000000 00:00 0        [heap]
...
35b180000-35b1820000 r-xp 00000000 08:02 135522    /usr/lib64/ld-2.15.so
35b1a1f000-35b1a20000 r--p 0001f000 08:02 135522    /usr/lib64/ld-2.15.so
35b1a20000-35b1a21000 rw-p 00020000 08:02 135522    /usr/lib64/ld-2.15.so
35b1a21000-35b1a22000 rw-p 00000000 00:00 0
35b1c00000-35b1dac000 r-xp 00000000 08:02 135870    /usr/lib64/libc-2.15.so
35b1dac000-35b1fac000 ---p 001ac000 08:02 135870    /usr/lib64/libc-2.15.so
35b1fac000-35b1fb0000 r--p 001ac000 08:02 135870    /usr/lib64/libc-2.15.so
35b1fb0000-35b1fb2000 rw-p 001b0000 08:02 135870    /usr/lib64/libc-2.15.so
```

- The currently mapped memory regions and their access permissions.
- The addresses shown in this file are Virtual memory address range
- We'll use all of VMAR to find pages status in **/proc/\$PID/pagemaps**

/proc/\$PID/pagemaps



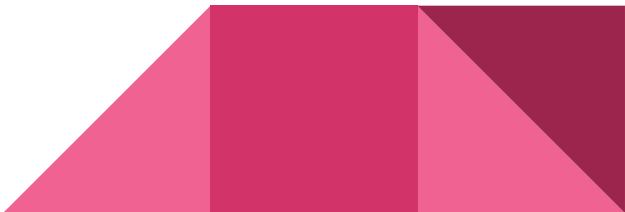
- Arrays of the page status. Each contained in 64-bit structure
- If the page is in the physical memory, the “Page Present” bit is 1 and it will contain **Page Frame Number (PFN)**
- At first, We'll use these PFN to determine page flag in **/proc/kpageflag**
- But now we must use **kernel mode...**

/proc/kpageflags

The flags are (from fs/proc/page.c, above kpageflags_read):

- 0. LOCKED
- 1. ERROR
- 2. REFERENCED
- 3. UPTODATE
- 4. DIRTY
- 5. LRU
- 6. ACTIVE
- 7. SLAB
- 8. WRITEBACK
- 9. RECLAIM
- 10. BUDDY
- 11. MMAP
- 12. ANON
- 13. SWAPCACHE
- 14. SWAPBACKED
- 15. COMPOUND_HEAD
- 16. COMPOUND_TAIL
- 17. HUGE
- 18. UNEVICTABLE
- 19. HWPOISON
- 20. NOPAGE
- 21. KSM
- 22. THP
- 23. BALLOON
- 24. ZERO_PAGE
- 25. IDLE

For SLAB and BUDDY allocation,
If the bit is 1, the method is used
to allocate the memory



Memory Allocation

Direct Memory Access (DMA) Pages - less than 16 MB

Normal addressable pages (NormalMem) - 16-896 MB

Dynamically mapped pages (HighMem) - more than 896 MB



Gathering Page Flag

In **kernel mode**, we can use sets of the following function, and structure element to get data:

1. pfn_to_page -- convert PFN to physical memory page detail.
2. page_zone -- get detail of memory zone used in the page.
3. (Optional) page_zonenum -- get the zone index used in the page
4. page->flags -- get flags contained in the page

```
154
155 #define pfn_to_page(pfn) ({
156     unsigned long __pfn = (pfn);
157     struct pglist_data *pgdat;
158     pgdat = __virt_to_node((unsigned long)pfn_to_virt(__pfn));
159     pgdat->node_mem_map + (__pfn - pgdat->node_start_pfn);
160 })
161 #define page_to_pfn(_page) ({
919
920 static inline struct zone *page_zone(const struct page *page)
921 {
922     return &NODE_DATA(page_to_nid(page))->node_zones[page_zonenum(page)];
923 }
924
```

page_mm.h

mm.h



Kernel Module

Kernel module is a new way to compile kernel by building process for external loadable modules that has been integrated into standard kernel build mechanism.

What we need to compile a basic kernel module :

- A Makefile to compile a module

```
obj-m += pfn2zone.o

all:
    make -C /lib/modules/$(shell uname -r)/build M=$(PWD) modules

clean:
    make -C /lib/modules/$(shell uname -r)/build M=$(PWD) clean
```

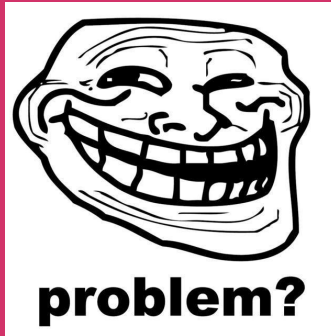
Inside of a Makefile file

Kernel Module

- In this case, we have a .C module file named pfn2zone.c
- In terminal, use command “make” to compile it.
- Then, use “insmod ./pfn2zone.ko [parameters]” to insert to kernel
- In case of there's a kernel module that we don't want it anymore use “rmmod modulename” to remove it from your kernel.

```
9  MODULE_LICENSE("GPL");
10 MODULE_AUTHOR("JoeJa");
11
12 //unsigned long long pfn = 0;
13
14 static unsigned long long pfn_array[3000];
15 static int pfnArrayCnt = 0;
16
17 module_param_array(pfn_array, ullong, &pfnArrayCnt, 0000);
18 MODULE_PARAM_DESC(pfn_array, "An array of Page Frame Number");
19
20 void getZone(unsigned long long pfn){
21     struct page *pageObj;
22     struct zone *pageZone;
23
24     //printk(KERN_INFO "NID: %d=>%d\n", pfn_valid(pfn), pfn_to_nid(pfn));
25     if(pfn_valid(pfn)){
26         pageObj = pfn_to_page(pfn);
27         pageZone = page_zone(pageObj);
28         // printk(KERN_INFO "Zone Num: %d\t", page_zonenum(pageObj));
29         // printk(KERN_INFO "Zone: %s\n", pageZone->name);
30         printk(KERN_INFO "pfn2zone_by_Joe 0x%llx 0x%lx %d %s\n", pfn, pageObj->flags, page_zonenum(pageObj), pageZone->name);
31     }else{
32         printk(KERN_INFO "pfn2zone_by_Joe 0x0 0x0 -1 Invalid\n");
33     }
34 }
```

Problem



Reading /proc/\$PID/pagemap

```
tam@tam-VirtualBox:~$ cat /proc/1254/pagemap
```

A large grid of small, repeating circular patterns, likely a decorative background or a placeholder for a large image. The patterns are arranged in a dense, uniform grid across the entire page. Each pattern consists of a central dot surrounded by concentric circles, creating a textured, woven appearance. The colors are muted, with shades of brown, tan, and grey, giving it a vintage or rustic feel.

Issue in the kernel mode

- Result must show output via printk -> results are in the “DMESG”
 - Solution: Create DMESG reader, add signature prefix before the real interested value. So we can use **grep** to filter out unwanted data rows from another programs.
 - Use * trick from sscanf: ignore the following format value.

```
osboxes@osboxes:~/os_memory$ gcc how2read_dmesg.c -o how2read
osboxes@osboxes:~/os_memory$ sudo rmmod pfn2zone
osboxes@osboxes:~/os_memory$ sudo insmod pfn2zone.ko pfn_array=0x110e8e,0x1160e5,0x100c17,0x100c14,0x112c
4d,0xd1065,0x109b40
osboxes@osboxes:~/os_memory$ sudo ./how2read
RAW: [14163.087697] pfn2zone_by_joe 0x112c4d 2 Normal
                    pfn=112c4d, zone ID=2, zone Name=Normal
RAW: [14163.087698] pfn2zone_by_joe 0xd1065 1 DMA32
                    pfn=d1065, zone ID=1, zone Name=DMA32
RAW: [14163.087698] pfn2zone_by_joe 0x109b40 2 Normal
                    pfn=109b40, zone ID=2, zone Name=Normal
osboxes@osboxes:~/os_memory$
```

```
osboxes@osboxes:~/os_memory$ sudo cat /dev/kmsg | grep pfn2zone_by_joe | tail -xd", cnt));
```

```
328
329 while (fgets(buf, sizeof(buf)-1, fp) != NULL) {
330     if(DEBUG) printf("RAW: %s\t\t", buf);
331     if(sscanf(
332         strchr(buf, ']'),
333         "%*s %*s %llx %llx %d %s",
334         &pfn,
335         &flags,
336         &zoneId,
337         zoneName
338     ) < 4){
339         continue; //ignore defect
340     }else{
341         currAddr = (MEMORY_ZONE_T*) malloc(sizeof(MEMORY_ZONE_T));
342         if(currAddr == NULL){
343             perror("Error while allocating address list: ");
344             exit(EXIT_FAILURE);
345         }
346
347         currAddr->pfn = pfn;
348         currAddr->flags = flags;
349         currAddr->zoneId = zoneId;
350         currAddr->name = (char*)malloc(sizeof(char)*strlen(zoneName));
351         strncpy(currAddr->name, zoneName, strlen(zoneName));
352
```

So
Kernel
:



Demo





Q&A

Reference

- <https://fivelinesofcode.blogspot.com/2014/03/how-to-translate-virtual-to-physical.html>
- <http://lxr.free-electrons.com/source/tools/vm/page-types.c>
- <http://www.tldp.org/LDP/lkmpg/2.6/html/lkmpg.html>
- https://www.gnu.org/software/libc/manual/html_node/Permission-Bits.html
- <https://stackoverflow.com/questions/5947286/how-can-linux-kernel-modules-be-loaded-from-c-code>
- <http://www.cplusplus.com/reference/cstdio/scanf/>

